

Python and Databases

GNBF5010 Lecture 11

2021R1

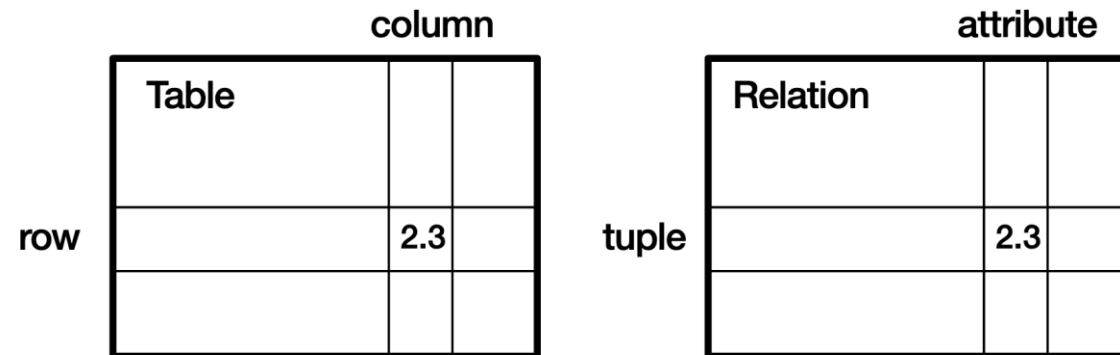
Overview

- Introduction to Databases and SQL
- Interacting with SQLite from Python

Introduction to Databases

What is a database

- An **organized** collection of **structured information**, or **data**, typically stored electronically in a computer system.
- Usually controlled by a database management system (**DBMS**)
 - e.g. MySQL, Microsoft SQL Server, Oracle Database, and **SQLite**.



- A relational database looks like a spreadsheet with multiple sheets (tables, or relations)

Database vs Dictionary

- Both map from *keys* to *values*.
- Both are designed to keep the **inserting and accessing** of data very **fast**, even for large amounts of data.
- But a database is on a **permanent storage**, and thus persists after program ends.
- A database can **store far more data** than a dictionary.
 - A python dictionary is limited to the size of the computer memory.

<https://www.py4e.com/html3/15-database>

Database vs Spreadsheet

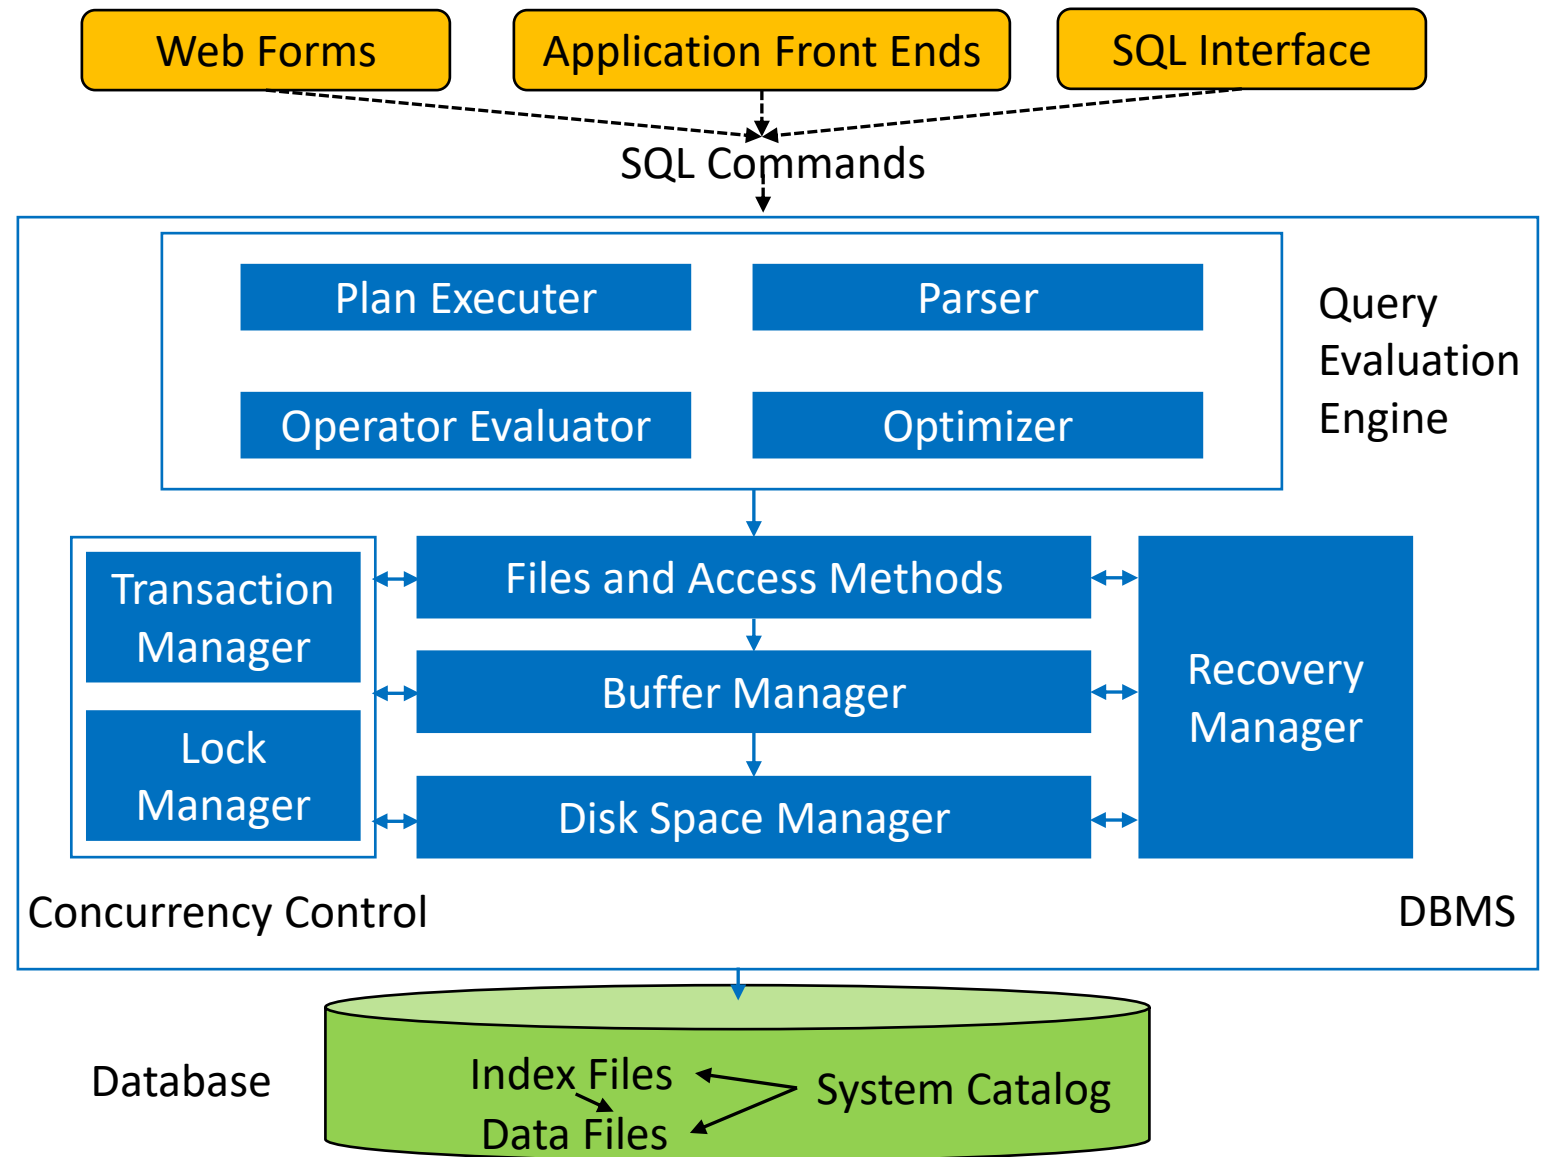
- Spreadsheets were originally designed for **one user** (or a few) who don't need to do a lot of complicated data manipulation.
- Databases are designed to hold much **larger** collections of organized information - **massive amounts**, sometimes.
- Databases allow **multiple users** at the same time to quickly and securely access and query the data **using highly complex logic** and **language**.

<https://www.oracle.com/hk/database/what-is-database/>

What is a database system

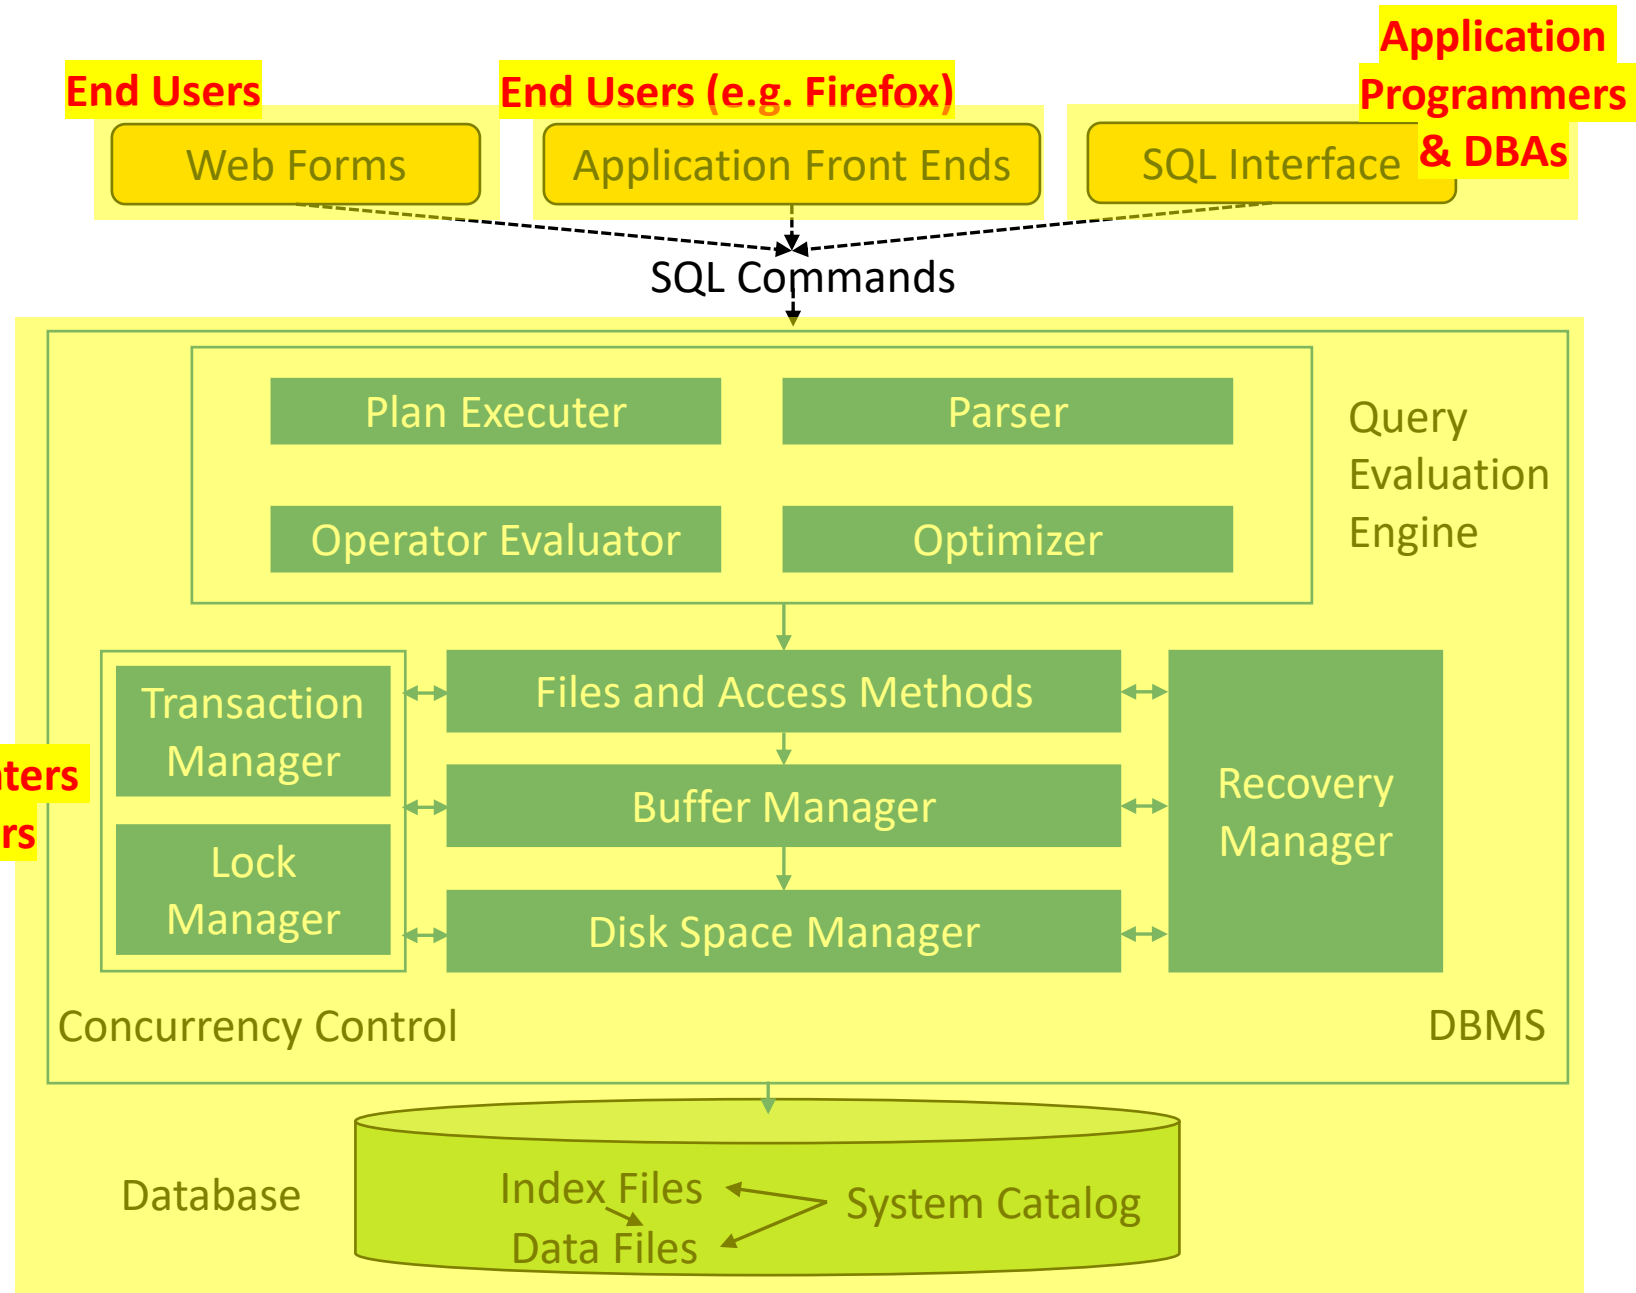
- A **database system** (or shortened to **database**)
= **Data + DBMS + Applications**
- The **data** can then be **easily** accessed, managed, modified, updated, controlled, and organized.
- Most databases use structured query language (**SQL**) for **writing** and **querying** data.

Architecture of a database system



People who work with databases

DBMS Implementers and Researchers



Common Database Systems

- Three major Database Management Systems (DBMS) in wide use
 - Oracle - Large, commercial, enterprise-scale, very very tweakable
 - MySQL - Simpler but very fast and scalable, commercially open source
 - SqlServer - Very nice, from Microsoft
- Many other smaller projects, free and open source
 - HSQL, **SQLite**, Postgres, ...

The SQLite DBMS www.sqlite.org

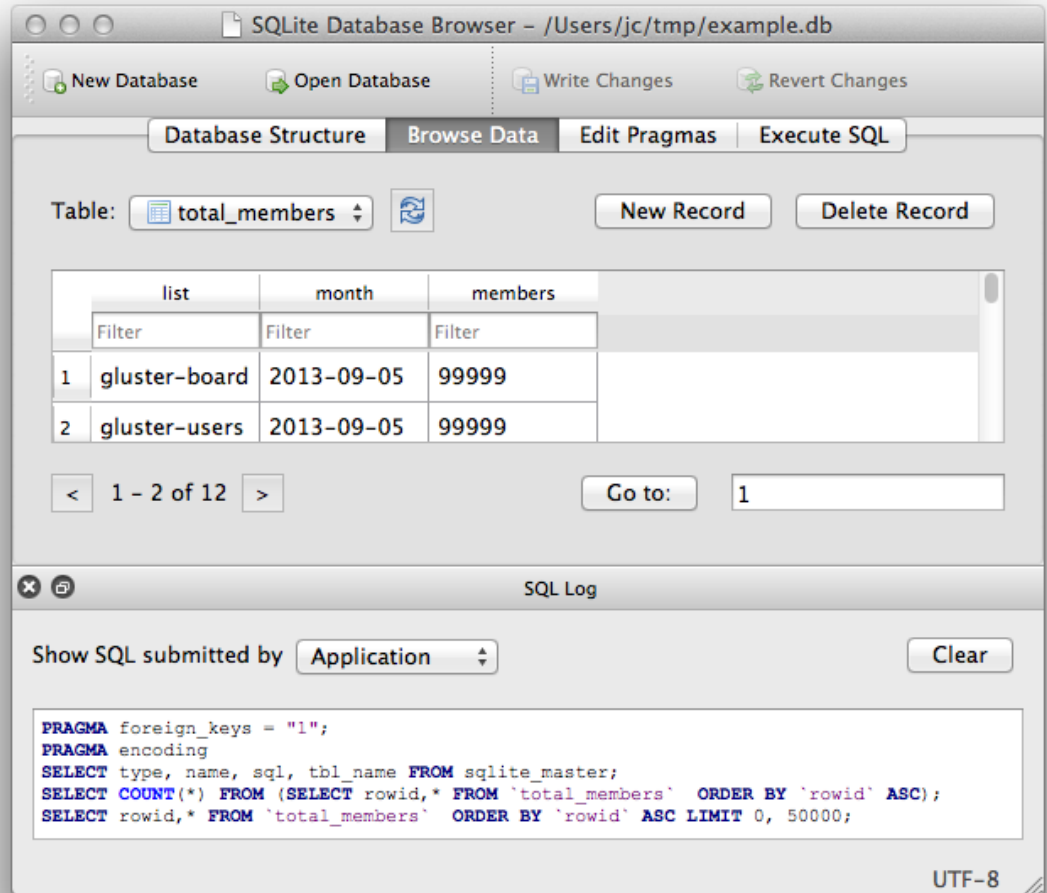
- **SQLite** is widely-used and is favourite among the developers for many reasons:
- Extremely light-weighted (not more than 500 KBs)
- Serverless – don't need any separate server for availing its services
- No complex setup
- The data can be saved in a single file – easy to transfer

SQLite is used by lots of software

The Symbian logo, featuring the word "symbian" in a bold, lowercase, sans-serif font. The letter "i" is stylized with a blue dot and a blue vertical bar.The McAfee logo, featuring the word "McAfee" in a bold, red, sans-serif font.The Microsoft logo, featuring the word "Microsoft" in a bold, black, sans-serif font.The Toshiba logo, featuring the word "TOSHIBA" in a bold, red, sans-serif font.The Google logo, featuring the word "Google" in its multi-colored, sans-serif font.

DB Browser for SQLite

www.sqlitebrowser.org



How to make a database

Structured Query Language (SQL)

The language we use to issue commands to the database

- Create data (a.k.a Insert)
- Retrieve data
- Update data
- Delete data

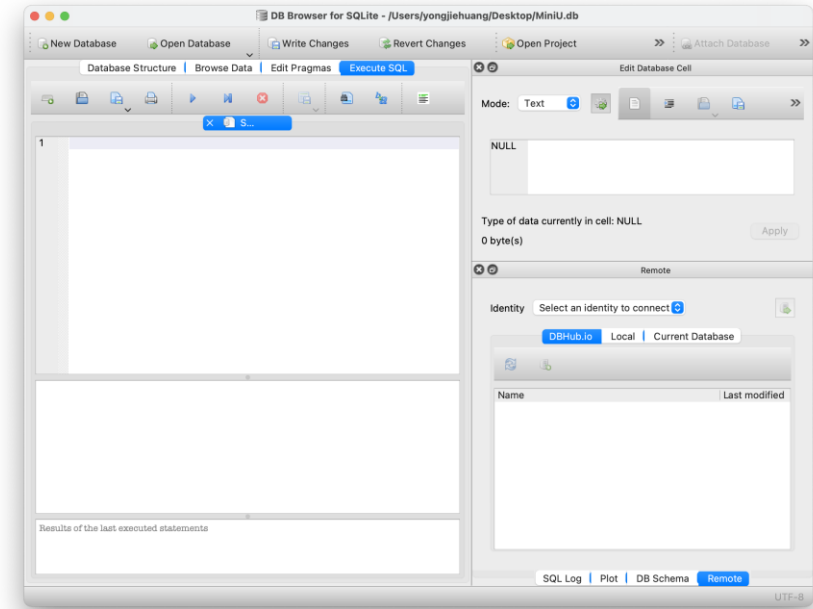
Create a database using SQLite

- In SQLite Browser, click "New Database" button.
- Alternatively, in command line, run
\$ **sqlite3** <db_file>
which will enter SQLite command line. To exit, type **.exit**
See more at <https://tool.oschina.net/uploads/apidocs/sqlite/sqlite.html>

Create tables

```
CREATE TABLE Students (  
    SSN      INTEGER PRIMARY KEY,  
    Name     TEXT,  
    Addr     TEXT  
);  
INSERT INTO Students Values(123, "smith", "main str");  
INSERT INTO Students Values(456, "jones", "QF ave");  
  
CREATE TABLE Classes (  
    Cid      TEXT PRIMARY KEY,  
    CName    TEXT,  
    Units    INTEGER  
);  
INSERT INTO Classes Values("15-413", "s.e.", 2);  
INSERT INTO Classes Values("15-412", "o.s.", 2);  
  
CREATE TABLE "Takes" (  
    "SSN"     INTEGER,  
    "Cid"     TEXT,  
    "Grade"   TEXT,  
    FOREIGN KEY("Cid") REFERENCES "Classes"("Cid"),  
    FOREIGN KEY("SSN") REFERENCES "Students"("SSN")  
);  
INSERT INTO Takes Values(123, "15-413", "A");  
INSERT INTO Takes Values(456, "15-413", "B");
```

lec11-create-miniU.sql



Execute commands in "Execute SQL" box in SQLite Browser; or
run from command line (`$ sqlite3 MiniU.db`)

Our Mini-U DB

Students		
<u>SSN</u>	Name	Addr
123	smith	main str
234	jones	QF ave

Classes		
<u>CId</u>	CName	Units
15-413	s.e.	2
15-412	o.s.	2

Takes		
<u>SSN</u>	<u>CId</u>	Grade
123	15-413	A
234	15-413	B

Insertions, Deletions and Updates

Insertions

```
insert into Students(ssn, name, address)  
values (123, 'smith', 'main')
```

OR...

```
insert into Students  
values (123, 'smith', 'main')
```

Deletions

- Delete the record of 'smith'

```
delete from Students  
where name='smith'
```

Be careful - it deletes ALL the 'smith's!

Updates

- Update the grade to 'A' for ssn=123 and course 15-415

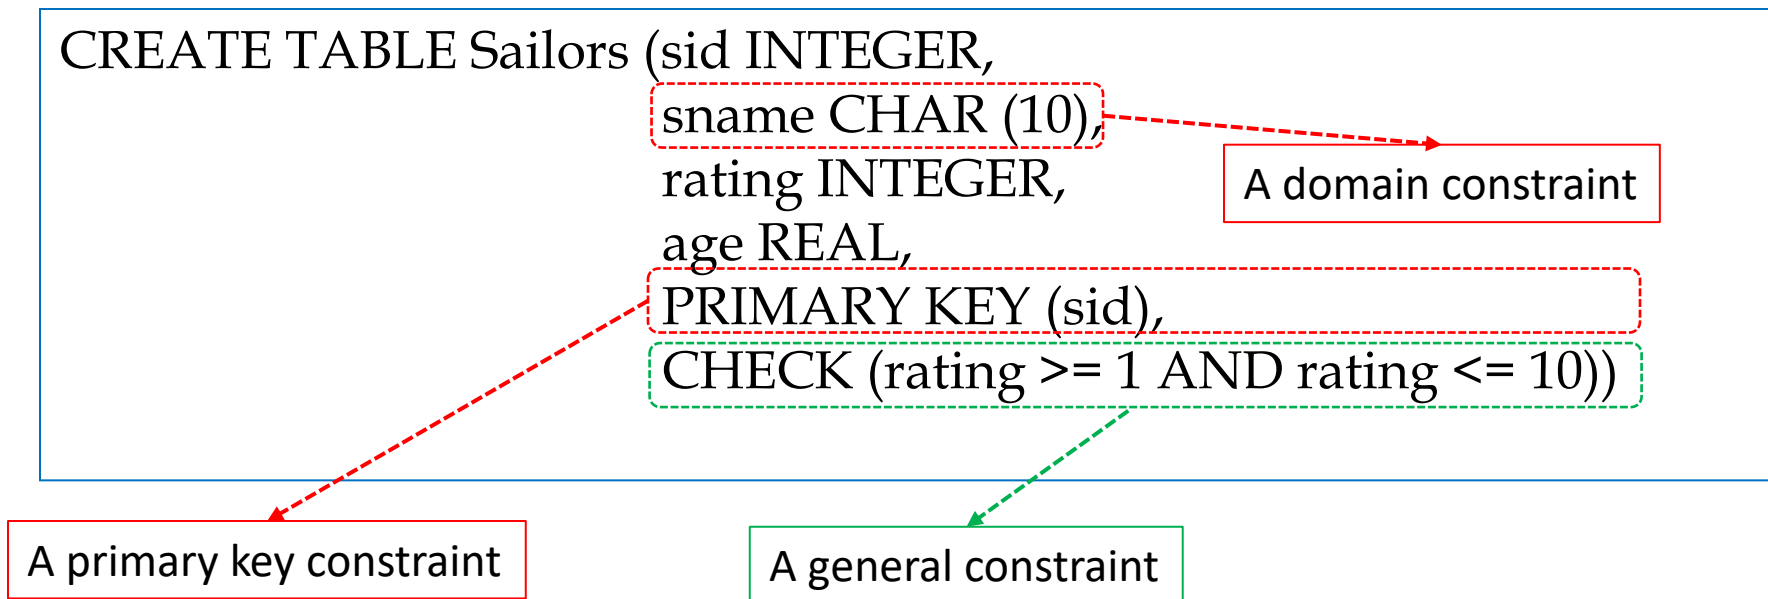
```
update takes  
set grade='A'  
where ssn = 123 and cid= '15-415'
```

Integrity Constraints- A Review

- An Integrity Constraint (IC) describes conditions that every *legal instance* of a relation must satisfy
- Inserts/deletes/updates that violate IC's are disallowed
- ICs can be used to:
 - Ensure application semantics (e.g., *sid* is a key)
 - Prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 20)

General Constraints Over a Single Table


- Complex constraints over a single table can be defined using **CHECK** *conditional-expression*



Basic SQL Queries

Basic SQL Queries

- The basic form of an SQL query is as follows:



Our Mini-U DB

Students		
<u>SSN</u>	Name	Addr
123	smith	main str
234	jones	QF ave

Classes		
<u>CId</u>	CName	Units
15-413	s.e.	2
15-412	o.s.	2

Takes		
<u>SSN</u>	<u>CId</u>	Grade
123	15-413	A
234	15-413	B

The WHERE Clause

- Find the ssn(s) of everybody called “smith”

Students		
<u>SSN</u>	Name	Addr
123	smith	main str
234	jones	QF ave

```
select ssn  
from students  
where name='smith'
```

The WHERE Clause

- Find ssn(s) of all “smith”s on “main”

Students		
<u>SSN</u>	Name	Addr
123	smith	main str
234	jones	QF ave

```
select ssn  
from students  
where addr='main' and  
      name = 'smith'
```

The WHERE Clause

- Boolean operators (**and, or, not**)
- Comparison operators (<, ≤, >, ≥, =, ≠)
- And more...

What About Strings?

- Find student ssn(s) who live on “main” (st or str or street – i.e., “main st” or “main str” or “main street”)

```
select ssn  
from students  
where addr like 'main%'
```

%: Variable-length do not care (i.e., stands for 0 or more arbitrary characters)

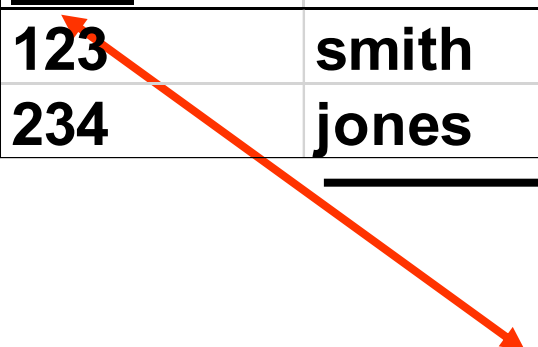
_: Single-character do not care (i.e., stands for any 1 character)

The FROM Clause

- Find the names of students taking 15-415

Students		
<u>SSN</u>	Name	Addr
123	smith	main str
234	jones	QF ave

Classes		
<u>CId</u>	CName	Units
15-413	s.e.	2
15-412	o.s.	2



Takes		
<u>SSN</u>	<u>CId</u>	Grade
123	15-413	A
234	15-413	B

The FROM Clause

- Find the names of students taking 15-415

```
select Name  
from STUDENTS, TAKES  
where ???
```


The FROM Clause

- Find the names of students taking 15-415

```
select Name  
from STUDENTS, TAKES  
where STUDENTS.ssn = TAKES.ssn  
        and TAKES.cid = '15-415'
```

Renaming: Tuple Variables

- Find the names of students taking 15-415

```
select Name  
from STUDENTS as S, TAKES as T  
where S.ssn = T.ssn  
        and T.cid = "15-415"
```

Optional!

Renaming: Self-Joins

- Find names and increments for the ratings of persons who have sailed two different boats on the same day

Sailors			
Sid	Sname	Rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0

Reserves		
Sid	Bid	Day
22	101	10/10/2013
22	102	10/10/2013

Note: Red arrows point from the text 'Sailor ID' to the 'Sid' column and 'Boat ID' to the 'Bid' column.

Renaming: Self-Joins

- Find names and increments for the ratings of persons who have sailed two different boats on **the same day**

Sailors			
Sid	Sname	Rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0

Reserves		
Sid	Bid	Day
22	101	10/10/2020
22	102	10/10/2020

```
select S.sname, S.rating+1 as rating
from Sailors S, Reserves R1, Reserves R2
where S.sid = R1.sid and S.sid = R2.sid
      and R1.day = R2.day and R1.bid != R2.bid
```

*R1 and R2 are different table
aliases for the same table.*

Renaming: Theta Joins

- Find course names with **more units than** 15-415

Classes		
<u>CId</u>	CName	Units
15-413	s.e.	2
15-412	o.s.	2

```
select c1.cname
from class as c1, class as c2
where c1.units > c2.units
      and c2.cid = '15-415'
```

Set Operations

Set Operations

- Find ssn(s) of students taking both 15-415 and 15-413

Takes		
<u>SSN</u>	<u>CId</u>	Grade
123	15-413	A
234	15-413	B

```
select ssn  
from takes  
where cid='15-415' and  
cid='15-413'
```

Set Operations

- Find ssn(s) of students taking both 15-415 and 15-413

Takes		
<u>SSN</u>	<u>Cid</u>	Grade
123	15-413	A
234	15-413	B

(select ssn from takes where cid="15-415")

intersect

(select ssn from takes where cid="15-413")

Other operations: **union , except**

Set Operations

- Find ssn(s) of students taking 15-415 or 15-413

Takes		
<u>SSN</u>	<u>Cid</u>	Grade
123	15-413	A
234	15-413	B

(select ssn from takes where cid="15-415")

union

(select ssn from takes where cid="15-413")

Set Operations

- Find ssn(s) of students taking 15-415 but not 15-413

Takes		
<u>SSN</u>	<u>Cid</u>	Grade
123	15-413	A
234	15-413	B

(select ssn from takes where cid="15-415")

except

(select ssn from takes where cid="15-413")

SQL Summary

```
INSERT INTO Users (name, email) VALUES ('Kristin', 'kf@umich.edu')
```

```
DELETE FROM Users WHERE email='ted@umich.edu'
```

```
UPDATE Users SET name="Charles" WHERE email='csev@umich.edu'
```

```
SELECT * FROM Users
```

```
SELECT * FROM Users WHERE email='csev@umich.edu'
```

```
SELECT * FROM Users ORDER BY email
```

See more examples at [py4e-ch15-handout.txt](#)

Interacting with SQLite from Python

Database API

- The `sqlite3` tool is primarily useful for extracting data from SQLite databases from a script, or quick exploration and interaction with a SQLite database.
- But for more involved tasks such as `loading numerous records` into a database or `executing complex queries` as part of data analysis, it's often preferable to interact with a SQLite database through an `API`.
- APIs allow you to `interact` with a database through an interface in your language of choice. We will take a quick look at Python's excellent API.

Python module - sqlite3

```
import sqlite3
```

This must be the first line of the program to allow Python to use the SQLite3 library.

```
with sqlite3.connect("company.db") as db:  
    cursor=db.cursor()
```

Connects to the company database. If no such database exists, it will create one. The file will be stored in the same folder as the program.

```
cursor.execute("""CREATE TABLE IF NOT EXISTS employees(  
    id integer PRIMARY KEY,  
    name text NOT NULL,  
    dept text NOT NULL,  
    salary integer);""")
```

Creates a table called employees which has four fields (id, name, dept and salary). It specifies the data type for each field, defines which field is the primary key and which fields cannot be left blank. The triple speech marks allow the code to be split over several lines to make it easier to read rather than having it all displayed in one line.

db.cursor()

```
cursor.execute("""INSERT INTO employees(id,name,dept,salary)
VALUES ("1","Bob","Sales","25000") """)
db.commit()
```

Inserts data into the employees table. The **db.commit()** line saves the changes.

```
newID = input("Enter ID number: ")
newname = input("Enter name: ")
newDept = input("Enter department: ")
newSalary = input("Enter salary: ")
cursor.execute("""INSERT INTO employees(id,name,dept,salary)
VALUES (?, ?, ?, ?) """, (newID,newName,newDept,newSalary))
db.commit()
```

Allows a user to enter new data which is then inserted into the table.

```
cursor.execute("SELECT * FROM employees")  
print(cursor.fetchall())
```

Displays all the data from the employees table.

```
cursor.execute("SELECT * FROM employees")  
for x in cursor.fetchall():  
    print(x)
```

Displays all the data from the employees table and displays each record on a separate line.



```
cursor.execute("SELECT * FROM employees ORDER BY name")  
for x in cursor.fetchall():  
    print(x)
```

Selects all the data from the employees table, sorted by name, and displays each record on a separate line.


```
cursor.execute("SELECT * FROM employees WHERE salary>20000")
```

Selects all the data from the employees table where the salary is over 20,000.

```
cursor.execute("SELECT * FROM employees WHERE dept='Sales'")
```

Selects all the data from the employees table where the department is "Sales".

```
cursor.execute("""SELECT employees.id,employees.name,dept.manager  
FROM employees,dept WHERE employees.dept=dept.dept  
AND employees.salary >20000""")
```

Selects the ID and name fields from the employees table and the manager field from the department table if the salary is over 20,000.

```
cursor.execute("SELECT id,name,salary FROM employees")
```

Selects the ID, name and salary fields from the employees table.

```
whichDept = input("Enter a department: ")  
cursor.execute("SELECT * FROM employees WHERE dept=?", [whichDept])  
for x in cursor.fetchall():  
    print(x)
```

Allows the user to type in a department and displays the records of all the employees in that department.

```
cursor.execute("""SELECT employees.id, employees.name, dept.manager  
    FROM employees, dept WHERE employees.dept=dept.dept""")
```

Selects the ID and name fields from the employees table and the manager field from the department table, using the dept fields to link the data. If you do not specify how the tables are linked, Python will assume every employee works in every department and you will not get the results you are expecting.

```
cursor.execute("UPDATE employees SET name = 'Tony' WHERE id=1")  
db.commit()
```

Updates the data in the table (overwriting the original data) to change the name to "Tony" for employee ID 1.

```
cursor.execute("DELETE employees WHERE id=1")
```

Connecting to SQLite databases and creating tables from Python

create_table.py

```
import sqlite3

# the filename of this SQLite database
db_filename = "variants.db"

# initialize database connection
conn = sqlite3.connect(db_filename) ❶

c = conn.cursor() ❷

table_def = """\ ❸
CREATE TABLE variants(
    id integer primary key,
    chrom test,
    start integer,
    end integer,
    strand text,
    rsid text);
"""

c.execute(table_def) ❹

conn.commit() ❺
conn.close() ❻
```

Loading data into a table from Python

load_variants.py

```
import sys
import sqlite3
from collections import OrderedDict

# the filename of this SQLite database
db_filename = "variants.db"

# initialize database connection
conn = sqlite3.connect(db_filename)
c = conn.cursor()

## Load Data
# columns (other than id, which is automatically incremented
tbl_cols = OrderedDict([("chrom", str), ("start", int), ❶
                        ("end", int), ("strand", str),
                        ("rsid", str)])

with open(sys.argv[1]) as input_file:
    for line in input_file:
        # split a tab-delimited line
        values = line.strip().split("\t")

        # pair each value with its column name
```

```

cols_values = zip(tbl_cols.keys(), values) ❷

# use the column name to lookup an appropriate function to coerce each
# value to the appropriate type
coerced_values = [tbl_cols[col](value) for col, value in cols_values] ❸

# create an empty list of placeholders
placeholders = ["?"] * len(tbl_cols) ❹

# create the query by joining column names and placeholders quotation
# marks into comma-separated strings
colnames = ", ".join(tbl_cols.keys())
placeholders = ", ".join(placeholders)
query = "INSERT INTO variants(%s) VALUES (%s);"%(colnames, placeholders)

# execute query
c.execute(query, coerced_values) ❺

conn.commit() # commit these inserts
conn.close()

```

Loading *.json
file into a table
from Python

roster.py

```
import json
import sqlite3

conn = sqlite3.connect('rosterdb.sqlite')
cur = conn.cursor()

# Do some setup
cur.executescript('''
DROP TABLE IF EXISTS User;
DROP TABLE IF EXISTS Member;
DROP TABLE IF EXISTS Course;

CREATE TABLE User (
    id      INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
    name    TEXT UNIQUE
);

CREATE TABLE Course (
    id      INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
    title   TEXT UNIQUE
);

CREATE TABLE Member (
    user_id    INTEGER,
    course_id  INTEGER,
    role       INTEGER,
    PRIMARY KEY (user_id, course_id)
)''')
```

```

fname = input('Enter file name: ')
if len(fname) < 1:
    fname = 'roster_data_sample.json'

# [
#   [ "Charley", "sill0", 1 ],
#   [ "Mea", "sill0", 0 ],

str_data = open(fname).read()
json_data = json.loads(str_data)

for entry in json_data:

    name = entry[0]
    title = entry[1]

    print((name, title))

    cur.execute('''INSERT OR IGNORE INTO User (name)
        VALUES ( ? )''', ( name, ) )
    cur.execute('SELECT id FROM User WHERE name = ? ', (name, ))
    user_id = cur.fetchone()[0]

    cur.execute('''INSERT OR IGNORE INTO Course (title)
        VALUES ( ? )''', ( title, ) )
    cur.execute('SELECT id FROM Course WHERE title = ? ', (title, ))
    course_id = cur.fetchone()[0]

    cur.execute('''INSERT OR REPLACE INTO Member
        (user_id, course_id) VALUES ( ?, ? )''',
        ( user_id, course_id ) )

conn.commit()

```

Work with the Cursor object

```
>>> import sqlite3
>>> conn = sqlite3.connect("variants.db")
>>> c = conn.cursor()
>>> statement = """\
... SELECT chrom, start, end FROM variants WHERE rsid IN ('rs12255372', 'rs333')
... """
>>> c.execute(statement)
<sqlite3.Cursor object at 0x10e249f80>
>>> c.fetchone()
(u'chr10', 114808901, 114808902)
>>> c.fetchone()
(u'chr3', 46414946, 46414978)
>>> c.fetchone() # nothing left
>>>
```


Dumping Databases

Dumping Databases

- useful to back up and duplicate databases
- also useful in sharing databases
 - In SQLite it's easier to simply share the database file though
 - but this isn't possible with other database engines like MySQL and PostgreSQL

Dumping database using SQLite

```
$ sqlite3 variants.db ".dump"  
PRAGMA foreign_keys=OFF;  
BEGIN TRANSACTION;  
CREATE TABLE variants(  
    id integer primary key,  
    chrom text,  
    start integer,  
    end integer,  
    strand text,  
    rsid text);  
INSERT INTO "variants" VALUES(1,'chr10',114808901,114808902,'+', 'rs12255372');  
INSERT INTO "variants" VALUES(2,'chr9',22125502,22125503,'+', 'rs1333049');  
INSERT INTO "variants" VALUES(3,'chr3',46414946,46414978,'+', 'rs333');  
INSERT INTO "variants" VALUES(4,'chr2',136608645,136608646,'-', 'rs4988235');  
COMMIT;
```

Use database dumps to create databases

```
$ sqlite3 variants.db ".dump" > dump.sql  
$ sqlite3 variants-duplicate.db < dump.sql
```

This series of commands dumps all tables in the *variants.db* database to a SQL file *dump.sql*. Then, this SQL file is loaded into a new empty database *variants-duplicate.db*, creating all tables and inserting all data in the original *variants.db* database.

Recommended readings

- See materials in the Blackboard.